

OpenCollab:
A Blockchain Based Protocol to Incentivize Open
Source Software Development and Governance

Yondon Fu

April 27, 2017

1 Introduction

In today's software dependent society, open source software is everywhere. Parties ranging from large technology corporations like Google to individual hobbyist developers use open source software as the building blocks of their own projects. Tools that a developer once had to build from scratch are now widely available for anyone to use for free on websites like Github. Not only can anyone easily use these software packages, but anyone can also freely access, inspect and alter the source code, tailoring it for his or her own specialized needs.

The implications of democratized access to quality software is wide ranging. The open source web framework Ruby on Rails not only powers popular applications such as Twitter and Github that millions of people rely on everyday, but it also made web application development accessible to a broader audience by abstracting away the details of composing together components such as HTTP request handling, database querying and templating. Given the importance of open source software, the task at hand for technologists is to figure out how to make open source software projects sustainable such that organizations and individuals in the future can continue to rely on them in the future.

The sustainability of an open source software project is tied to project health and support. Project health is determined by how actively and adequately project developers communicate with users such that the project addresses the needs of the community. Project support is determined by the availability of financial and technical resources to develop a project[6]. A sustainable open source software project needs to be both healthy and supported.

A healthy and supported project optimizes the use of developer time and attention, the scarcest resources in open source software projects. Communication between developers and users in a healthy project informs developers of community needs and issues to potentially focus their time and attention on. Availability of financial and technical resources in a supported project ensures developers are free to allocate their all of their time and attention on project issues. Consequently, in a healthy and supported project, developers can properly allocate their time and attention according to community needs.

If project developers fail to properly allocate their time and attention, users might leave a project in search of alternatives that better suit their needs. Canonical, the company behind the Linux operating system Ubuntu, created fragmentation in the Linux community when it shipped a new version of Ubuntu with the Unity interface rather than the standard GNOME interface[4]. Canonical's failure to properly poll for user opinion and allocate developer time and attention accordingly ultimately hurt the Ubuntu project.

1.1 Bounty Systems

A proposed solution to open source software sustainability is a bounty system for project issues such as the one operated by the website Bountysource[2]. In these systems, users can attach monetary bounties to project issues that are rewarded to contributors that successfully resolve issues. While these bounty systems

may help projects attract more contributors, they do not take into account the incentives of maintainers. The work done by maintainers to review and merge in code is just as crucial as the work done by contributors. Consequently, bounty systems only partially help with project support.

Furthermore, bounty systems do not necessarily help with project health. Although in some cases, multiple users placing bounties on an issue might signal the importance the community places on that particular issue, it is also possible for malicious actors to place large bounties on issues that would negatively impact project quality. Such a possibility can place a burden on developers of filtering signal from noise and also introduces the possibility of collusion - a contributor might share a large bounty if a maintainer agrees to merge it into the codebase even if the contributed code is of poor quality. As a result, bounty systems can actually hamper communication between developers and users leading to unmet community needs. Adding any form of financial compensation to open source software projects needs to take into align the incentives of all parties involved or else perverse incentives might arise leading to malicious behavior that harms the quality of the project.

Lastly, bounty systems operated by websites like Bountysource rely on a centralized entity to facilitate transactions. This reliance on a centralized entity not only results in a central point of failure, but can actually be more costly for users. For example, although users can freely transact within the bounty system, Bountysource charges a 10% withdrawal fee if a user wants to cash out. As a result, users choose between giving up a portion of their monetary rewards and giving up the numerous opportunities to use their monetary rewards for their own benefit outside the bounty system. This withdrawal fee discourages users from leaving the system which benefits Bountysource, but harms users.

Cryptocurrencies, blockchains and smart contracts introduce the possibility of building protocols with embedded economic incentives. These protocols obviate the need for a centralized entity to facilitate transactions within a distributed network of users.

1.2 Contributions

The primary contributions of this thesis are the following:

- A command line tool that enables a decentralized Git workflow for developing open source software without relying on a centralized service like Github (Section 3).
- A proof-of-concept blockchain based protocol to incentivize open source software development and governance (Section 4).

The motivation behind these contributions is to push the discussion on how to improve open source software sustainability. In particular, these contributions are attempts to answer the following questions relating to open source software sustainability.

- How can developers poll for user opinion on issue prioritization for a project?
- How can a project attract regular contributors?
- How can maintainers be incentivized to carefully review and merge pull requests such that the quality of a project is upheld?
- The set of maintainers for a project should be able to change over time. However, maintainers might be reluctant to leave a project in fear of creating a leadership vacuum. How can incentives be structured such that people want to become maintainers?

A detailed description of these contributions can be found in Sections [3](#) and [4](#).

2 Background

Despite only recently receiving more mainstream attention among developers and business people, cryptocurrencies and blockchains have a long history. The technical foundations of cryptocurrencies and blockchains evolved from past work by computer science researchers. Two of the most important areas of research are in electronic cash systems and digital time-stamping.

2.1 Electronic Cash Systems

The advent of the Internet and the mainstream adoption of computing devices allowed people to freely transfer data between each other. A logical extension of online data transfer is online cash transfer. However, creating an electronic cash system comes with a number challenges. Similar to distributed systems, a useable electronic system also requires the ACID properties[5]:

- Atomicity: a transaction either occurs completely or does not occur at all
- Consistency: relevant parties for a transaction agree on the critical facts of exchange
- Isolation: transactions do not interfere with each other
- Durability: if one or more computers crash, the system should be able to recover to the last consistent state

David Chaum attempted to design a system with some of these properties with his DigiCash project. In DigiCash, a bank issues and verifies electronic tokens that network users can exchange with one another[5]. One of the flaws of DigiCash is that it can only offer durable transactions in exchange for a loss of anonymity. Network users must present tokens to the bank for verification or be at risk of double spending attacks since electronic messages can be duplicated. Additionally, this reliance on the bank creates a bottleneck for system throughput and a central point of failure. In the scenario where an attacker gains access to the bank's private key, the attacker can create counterfeit tokens that are indistinguishable from valid tokens at will.

2.2 Digital Time-Stamping

The Internet and mainstream adoption of computing devices allowed for the widespread digitization of all types of documents. While the digitization of documents provided many benefits, it also came with the problem of how to certify the existence and time of creation or change of a digital document.

In 1991, Haber and Stornetta presented a time-stamping method for digital documents that consisted of certificates cryptographically signed by a time-stamping service. The certificates contain the hash of the document as well as linking information from a previous certificate which includes a hash of the

previous certificate's linking information[7]. The result is a hash linked chain of certificates that prevents the faking of time-stamps.

Bayer, Haber and Stornetta extended this time-stamping method using merkle trees. In the original time-stamping method, verification of a document time-stamp can require at most N steps by following the chain links to a time-stamp certificate that is trustworthy[1]. Instead of linking N hashes of documents, the hash values can be stored in a merkle tree. Participants can record the hashes of their own documents and the sibling hash values along the path from the document hash to the root of the merkle tree. Consequently, verification can be done in at most $\lg N$ steps by presenting the document hash and the $\lg N$ hashes on the path to the root. This modified time-stamping approach reduces storage requirements and verification time.

2.3 Blockchains

Blockchains combine learnings from previous electronic cash systems like Chaum's DigiCash and digital time-stamping methods. A blockchain is distributed ledger that is not controlled or managed by a central entity. The ledger is powered by a network of connected computers that use a consensus mechanism to reach agreement over shared data[9].

2.4 Cryptographic Primitives and Data Structures

2.5 Bitcoin

2.6 Ethereum

3 Decentralized Git Workflow

3.1 Mango

3.2 Extensions to Mango

4 OpenCollab Protocol

4.1 Protocol Roles

- **Curators:** curate project issues by staking tokens.
- **Contributors:** open pull requests to resolve issues by staking tokens.
- **Maintainers:** review and merge pull requests for issues by staking tokens.

4.2 OpenCollab Token

The OpenCollab token (OCT) powers the OpenCollab protocol. The value offered by the token is influence over an open source software project. Furthermore, the token serves the following purposes in the protocol:

- Used in a staking mechanism for issue curation. Curators stake tokens to signal the importance they place on an issue.
- Used in a staking mechanism for opening pull requests. Contributors stake a certain number of tokens when opening a pull request. If a contributor's pull request is closed without being merged in to the project, the contributor's staked tokens are destroyed. The possibility of losing staked tokens discourages contributors from opening pull requests unless they are confident about the quality of their contributions.
- Used in a staking mechanism for merging pull requests. Maintainers stake a certain number of tokens when they initiate a merge. Before a merge is finalized, a token holder can challenge a maintainer's merge to start a voting round. If token holders decide to veto a maintainer's merge, the maintainer's staked tokens are destroyed. The possibility of losing staked tokens discourages maintainers from merging pull requests that do not benefit a project. The challenge and voting process for a merge is described in more detail in [Section 4.5](#).
- Used in deposits for token holders that choose to participate as voters in protocol governance.

An initial allocation of tokens will be distributed so that the various protocol roles can be fulfilled by token holders. A project creator can initialize a Mango repository and mint a certain amount of tokens for the initial allocation. The initial allocation might be done using a token crowdsale or by disbursement at the discretion of the project creator.

4.3 Curating Issues

Curators stake a number of tokens to an issue to signal the importance that they place on the issue. Since curators lock up their tokens for a period of time when they stake tokens to an issue, they have limited curation power. Curators

exchange either funds for tokens by purchasing them on the secondary market or work for tokens by providing work to the project as a contributor or maintainer. Furthermore, since curators take on the risk of a fall in token value, they have skin in the game[8]. If curators signal importance for bad issues, developers poorly allocate their time and attention. If developers properly allocate their time and attention on resolving issues that would increase the quality of a project, the community might lose interest and less people would desire influence over the project leading to a fall in token value. Consequently, curators have an incentive to signal importance for issues that accurately reflect the needs of the community. As well curated issues are resolved, the value of the token would increase thereby benefiting curators.

4.4 Opening Pull Requests

Contributors open pull requests by staking CONTRIBUTORSTAKE tokens, where CONTRIBUTORSTAKE is a repository parameter.

If a contributor’s pull request is successfully merged by a maintainer, the contributor receives a portion of the issue’s token reward. The portion can be calculated as $\text{REWARD} - (\text{REWARD} * \text{MAINTAINERPERCENTAGE})$.

If a contributor’s pull request is closed without being merged into the project, the contributor’s CONTRIBUTORSTAKE staked tokens are destroyed.

4.5 Merging Pull Requests

Maintainers merge pull requests by staking MAINTAINERSTAKE tokens, where MAINTAINERSTAKE is a repository parameter.

If a maintainer wants to merge a pull request, it calls INITMERGEPULLREQUEST(ID) to signal an intent to merge a particular pull request and starts a challenge period. During this period, any token holder can challenge the maintainer by calling CHALLENGE() and staking CHALLENGERSTAKE tokens.

If a maintainer is not challenged during the challenge period, he can call MERGEPULLREQUEST(ID). The maintainer receives a portion of the issues’s token reward which can be calculated as $\text{REWARD} * \text{MAINTAINERPERCENTAGE}$.

If a maintainer is challenged during the challenge period, a voting period begins. Voting takes place using a two step commit and reveal protocol first formalized by Brassard, Chaum and Crepeau[3]. During the commit step, voters with a minimum VOTERDEPOSIT deposit in the smart contract vote to uphold or veto a maintainer’s merge by calling COMMITVOTE(HASH) with the cryptographic hash of their vote and a secret phrase. A vote to uphold is a 0 and a vote to veto is a 1. The secret phrase can be any random string only known to the voter. The value of the vote is secure from an attacker as long as only the voter knows the secret phrase used when generating the hash. We use the KECCAK256 hashing function since it is used by Ethereum.

During the reveal step, voters reveal the values of their votes by submitting the concatenation of their vote and secret phrase used in the commit step by calling REVEALVOTE(VOTE). The smart contract verifies that the submitted

vote corresponds with the committed hash and tallies up votes as voters reveal them. Finally, anyone can call `VOTERESULT()` which compares the number of uphold and veto votes. The value that receives the majority of vote ($\geq 50\%$) wins. Voters on the losing side of the vote are penalized such that `VOTER-PENALTYPERCENTAGE` is deducted from their deposits. Voters on the winning side of the vote are rewarded such that `VOTERREWARDPERCENTAGE` is added to their deposits.

If a maintainer’s merge decision is upheld, the maintainer is able to call `MERGE_PULLREQUEST(ID)` to finalize the merge. The challenger’s `CHALLENGER-STAKE` staked tokens are destroyed and the maintainer can claim his portion of the issue token reward.

If a maintainer’s merge decision is vetoed, the challenger’s staked tokens are returned and the maintainer’s `MAINTAINERSTAKE` tokens are destroyed and is removed from the maintainer set for the repository. Consequently, the former maintainer would not only lose the staked tokens, but also the economic value of future issue token rewards. The possibility of losing tokens and maintainer status serves to encourage maintainer to only merge pull requests that ensure the quality of the project.

During a voting period, a malicious maintainer might attempt to bribe voters to vote to uphold the merge decision. However, given N voters, a maintainer would have to bribe $N/2$ voters to ensure the merge decision is upheld. Since a two step commit and reveal protocol is used for voting, voters cannot know with certainty that other voters will accept the maintainer’s bribe. Furthermore, since a voter will be penalized if he is on the losing side of the vote and rewarded if he is on the winning side of the vote, the bribe would need to be greater than both the possible reward and penalty. An economically rational maintainer would not try to bribe voters to finalize a merge as long as the cost of the bribe multiplied across $N/2$ voters is greater than the issue token reward a maintainer stands to gain from merging in the pull request.

5 Conclusion

References

- [1] Dave Bayer, Stuart Haber., and W. Scott Stornetta.
“Improving the Efficiency and Reliability of Digital Time-Stamping”.
In: *Sequences II: Methods in Communication, Security and Computer Science*. 1993, pp. 329–334.
- [2] *Bounty Source: Support for Open-Source Software*.
<https://www.bountysource.com/>. Accessed: 2017-04-24.
- [3] Gilles Brassard, David Chaum, and Claude Crepeau.
“Minimum Disclosure Proofs of Knowledge”.
In: *Journal of Computer and System Sciences* 37 (1988), pp. 156–189.
- [4] Jon Brodtkin.
Ubuntu Unity is dead: Desktop will switch back to GNOME next year.
<https://arstechnica.co.uk/information-technology/2017/04/ubuntu-unity-is-dead-back-to-gnome/>.
2017.
- [5] L. Jean Camp, Marvin Sirbu, and J.D. Tygar.
“Token and Notational Money in Electronic Commerce”.
In: *Proceedings of the 1st USENIX Workshop on Electronic Commerce*.
1995, pp. 1–12.
- [6] Nadia Eghbal. *What success really looks like in open source*.
<https://medium.com/@nayafia/what-success-really-looks-like-in-open-source-2dd1facaf91c>.
2016.
- [7] Stuart Haber and W. Scott Stornetta.
“How to time-stamp a digital document”.
In: *Journal of Cryptology* 3.2 (1991), pp. 9–111.
- [8] Nassim N. Taleb and Constantine Sandis.
“The Skin In The Game Heuristic for Protection Against Tail Events”.
In: *Review of Behavioral Economics* 1 (2014), pp. 1–21.
- [9] Peter Van Valkenburgh. *What is "Blockchain" anyway?*
<https://coincenter.org/entry/what-is-blockchain-anyway>. 2017.